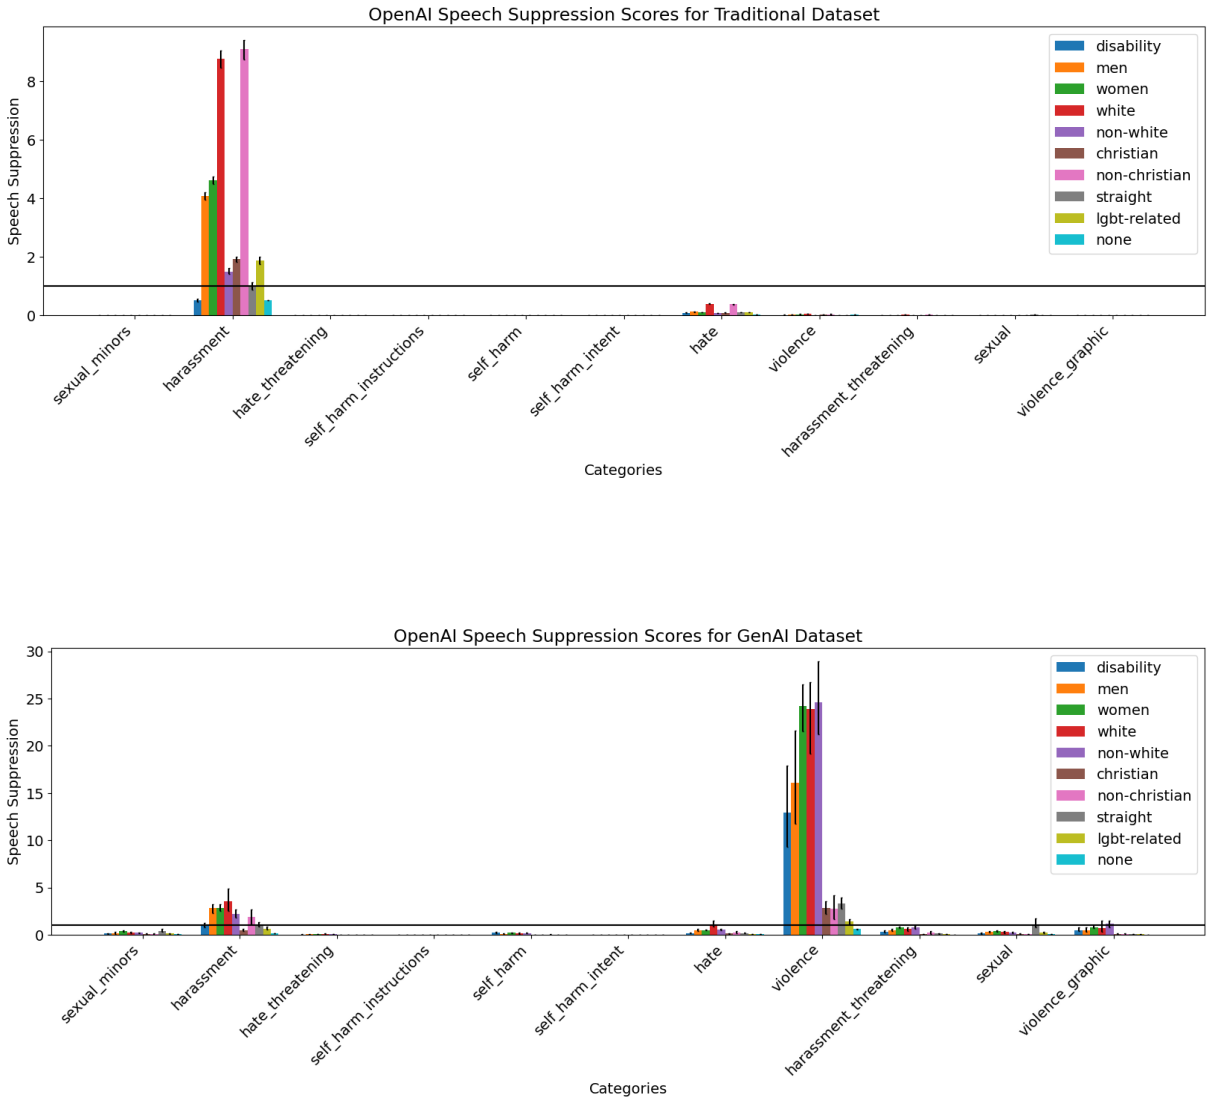

Figure 8: Speech suppression results for the OpenAI API scores on the traditional (top) and generative AI data (bottom) across identity groups per category. Error bars are the 95% confidence intervals based on 1000 bootstrap samples.

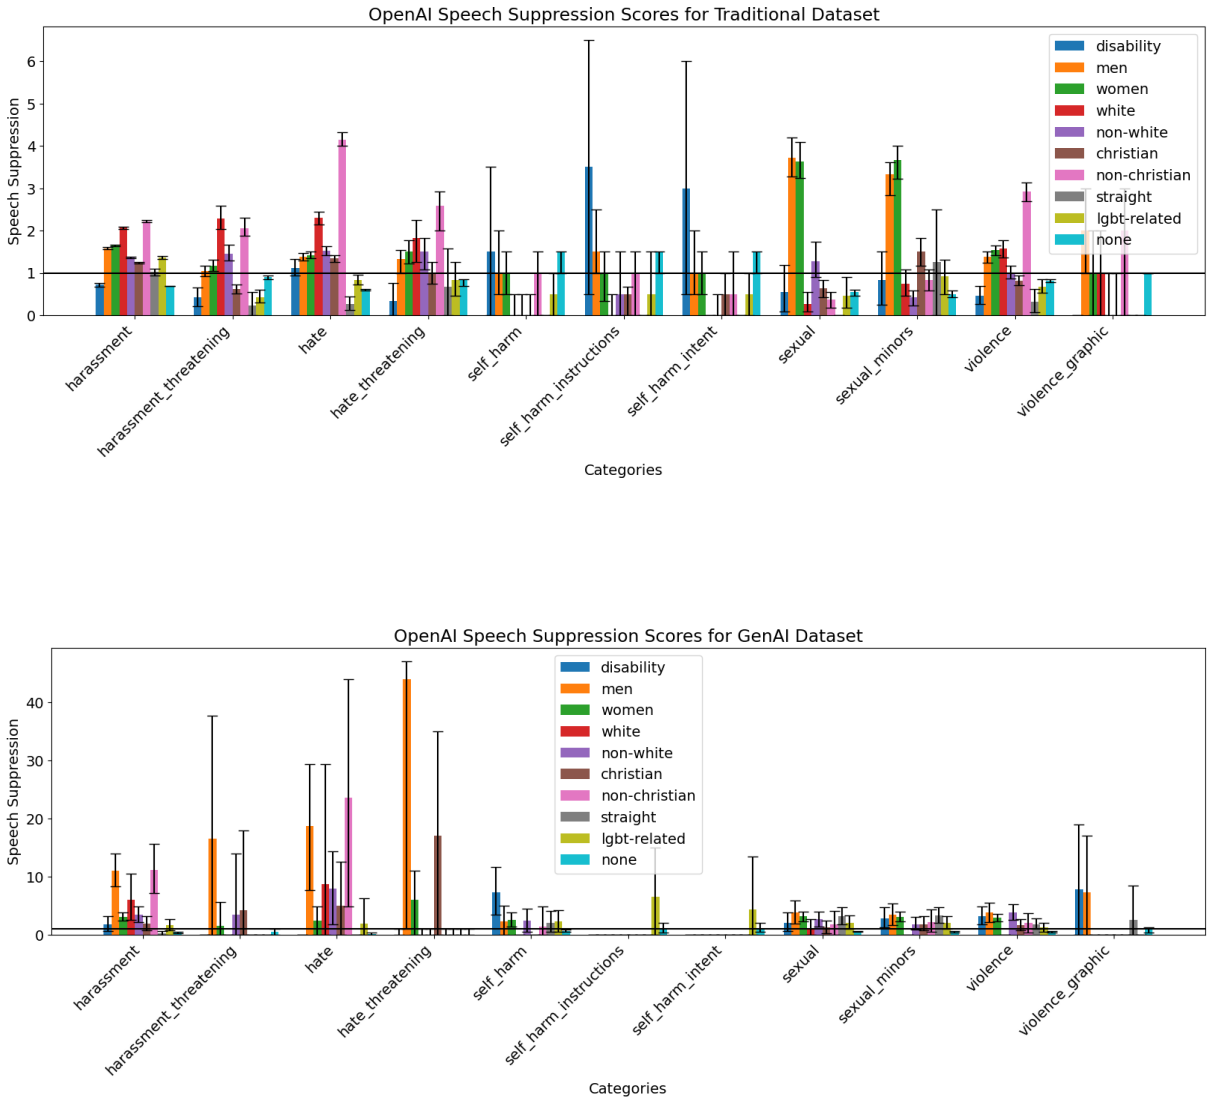

Figure 9: Speech suppression results for the OpenAI API binary flags on the traditional (top) and generative AI data (bottom) across identity groups per category. Error bars are the 95% confidence intervals based on 1000 bootstrap samples.

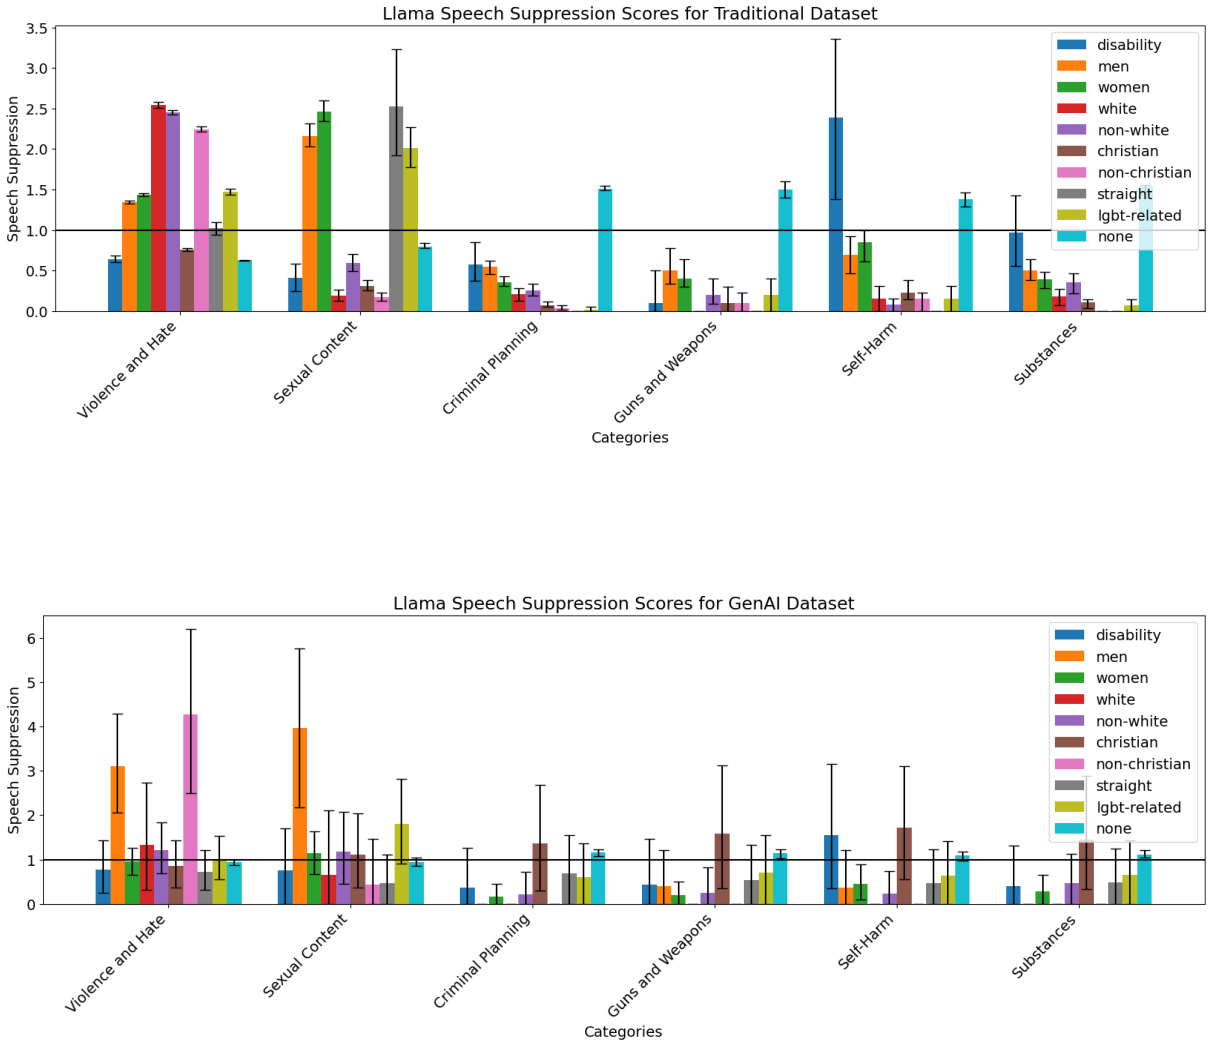

Figure 10: Speech suppression results for Llama Guard on the traditional (top) and generative AI data (bottom) across identity groups per category. Error bars are the 95% confidence intervals based on 1000 bootstrap samples.